

Advanced quantitative text analysis (2023W)

Petro Tolochko, Fabienne Lind

Day 3



Contents (smaller changes possible)

Day	Session 1	Session 2
1	Text as data, Text representation	Feature Engineering
2	Concepts & Data	Dictionaries
3	Supervised machine learning	Unsupervised machine learning
4	Neural network models, transformers	Using LLMs
5	Multilingual text analysis, Wrap-up	Project talks

Supervised classification

Day 3 Session 1

Types of machine learning

1. Supervised
 - An outcome variable is defined
 - Focus is on prediction
 2. Unsupervised
 - No outcome variable has been defined
 - Focus is on patterns
-

Types of machine learning

1. Supervised
 - An outcome variable is defined
 - Focus is on prediction
 2. Unsupervised
 - No outcome variable has been defined
 - Focus is on patterns
-

Supervised machine learning

Objectives:

- Classification (for categorical variables)
 - E.g.: classify documents into pre existing categories
 - Regression (for continuous variables)
-

Supervised machine learning

Objectives:

- **Classification (for categorical variables)**
 - **E.g.: classify documents into pre existing categories**
 - Regression (for continuous variables)
-

Steps

1. Create a labeled data set
2. Classify documents with supervised learning algorithm
3. Check performance
4. Using the measures

1. Create a labeled data set

How:

- Human coders annotate parts of the corpus (see also slides in session on dictionaries)
- Found data (e.g., self-reported profession in users' profile)

Considerations:

- Sampling should be representative for the corpus (e.g., Random, Stratified sample e.g., across time and source)
- Quality of human coding matters (Assess the intercoder reliability)
- Number of documents

1. Create a labeled data set

Number of documents

- the higher the number of categories and the lower the reliability of the coders, the higher the number of documents (Barberá et al., 2021)
- increase the sizes of manually coded validation dataset as large as possible, preferably to more than $N = 1,300$ (i.e., more than 1% of all data to be examined) assuming acceptable reliability (equal to or higher than .7) (Song et al., 2021)

Table 2. Simulation input parameters.

Factors	Input Parameters
N of human coders	2 (minimum), 5 (intermediate), & 10 (large manual coding)
Intercoder reliability	0.5 (low), 0.7 (acceptable), & 0.9 (high levels of reliability)
N of validation data	600 (0.5%), 1300 (1%), 6500 (5%), & 13000 (10%) of total data
Sampling variability	Random sample vs. nonrandom (biased) subset for validation
Coding per entry	Sole coding vs. duplicated coding for each entry

Song et al., 2021, p. 555

Song, H., Tolochko, P., Eberl, J. M., Eisele, O., Greussing, E., Heidenreich, T., ... & Boomgaarden, H. G. (2020). In validations we trust? The impact of imperfect human annotations as a gold standard on the quality of validation of automated content analysis. *Political Communication*, 37(4), 550-572.

Barberá, P., Boydston, A. E., Linn, S., McMahon, R., & Nagler, J. (2021). Automated text classification of news articles: A practical guide. *Political Analysis*, 29(1), 19-42.

1. Create a labeled data set

Split labeled data in training data, test data, validation set

Training data

- The subset that is used to learn the model parameters

Test data

- Another subset used to evaluate the model's predictive quality
- Not used for learning!

Validation data

- Only used to evaluate in the end, no further optimization allowed
-

2. Classify documents with supervised learning

Classifier learns the mapping between features and the labels in the training set

- We define a model $f(Y)=g(X)$
- And apply a learning algorithm to establish which features in X (features extracted from the training documents) matter to recover Y (i.e, the labels of the training documents)
- We fit the model

2. Classify documents with supervised learning

Considerations:

- Feature representation (Bag of words representation or embeddings)
- Feature selection (remove irrelevant features)
- Classifier selection
 - E.g., Naive Bayes, SVM, KNN, or ensemble methods

3. Check performance

The fitted model (the trained classifier) is applied to a held-out test set (which is a part of the labeled set but was not used for training the model).

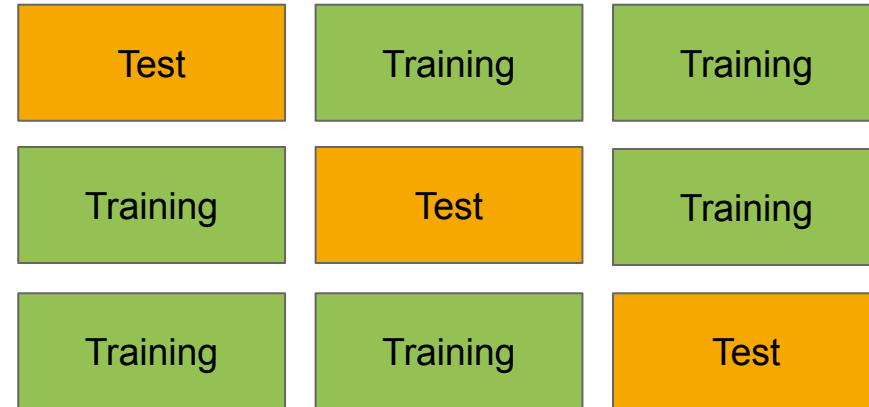
Considerations:

- Danger of overfitting (focus on features that work well with training set but do not generalize)
 - Solutions: cross-validation
- Performance metric (i.e., recall, precision)

3. Check performance

k-fold cross-validation

- We randomly split the data into k sets (“folds”) of roughly equal size
- Each set is hold out once as test set, while training on the remaining sets
- The problem of a lucky split is reduced



3. Check performance

Performance metrics

Confusion matrix

Classification (algorithm)	Actual label	
	Negative	Positive
Negative	True negative	False negative
Positive	False negative	True positive

$$Accuracy = \frac{True\ Negative + True\ Positive}{True\ Negative + True\ Positive + False\ Negative + True\ Positive}$$

$$Precision_{positive} = \frac{True\ Positive}{True\ Positive + False\ Positive}$$

$$Recall_{positive} = \frac{True\ Positive}{True\ Positive + False\ Negatives}$$

3. Check performance

Check convergent validity (Adcock & Collier, 2001)

- Compare the measures with other established measures, e.g.,
 - Use trained model to classify texts (open-ended answers relationship uncertainty as described by participants of an online survey) and compare it with the self-assessment in response to a closed question, see Pilny et al. 2019

Adcock, R., & Collier, D. (2001). Measurement validity: A shared standard for qualitative and quantitative research. *American political science review*, 95(3), 529-546.

Pilny, A., McAninch, K., Slone, A., & Moore, K. (2019). Using supervised machine learning in automated content analysis: An example using relational uncertainty. *Communication Methods and Measures*, 13(4), 287-304.

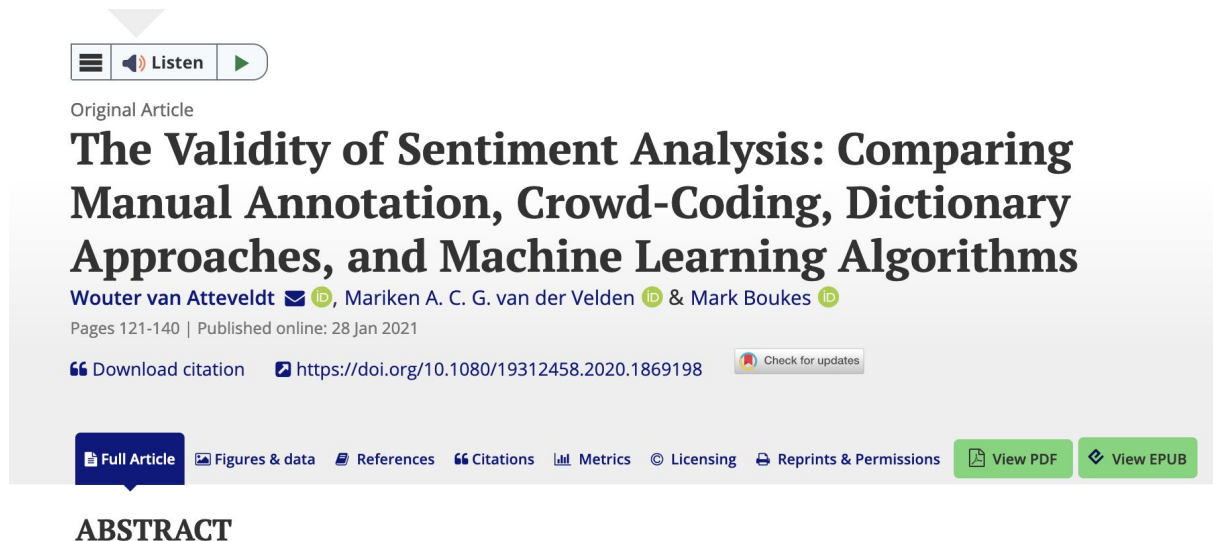
4. Using the measures

The classifier is applied to all documents in the corpus

Dictionary vs. supervised machine learning

Category: sentiment




Result: machine learning
significantly outperformed
dictionary coding






The screenshot shows the top portion of a research article page. At the top, there is a 'Listen' button with a speaker icon. Below it, the text 'Original Article' is displayed. The main title of the article is 'The Validity of Sentiment Analysis: Comparing Manual Annotation, Crowd-Coding, Dictionary Approaches, and Machine Learning Algorithms'. The authors listed are Wouter van Atteveldt, Mariken A. C. G. van der Velden, and Mark Boukes. Below the authors, it states 'Pages 121-140 | Published online: 28 Jan 2021'. There are links for 'Download citation' and a DOI link 'https://doi.org/10.1080/19312458.2020.1869198'. A 'Check for updates' button is also present. At the bottom of the article preview, there is a navigation bar with buttons for 'Full Article', 'Figures & data', 'References', 'Citations', 'Metrics', 'Licensing', 'Reprints & Permissions', 'View PDF', and 'View EPUB'. Below this navigation bar, the word 'ABSTRACT' is visible.










Original Article

The Validity of Sentiment Analysis: Comparing Manual Annotation, Crowd-Coding, Dictionary Approaches, and Machine Learning Algorithms

Wouter van Atteveldt , Mariken A. C. G. van der Velden  & Mark Boukes 

Pages 121-140 | Published online: 28 Jan 2021

 Download citation  <https://doi.org/10.1080/19312458.2020.1869198>  Check for updates

 Full Article  Figures & data  References  Citations  Metrics  Licensing  Reprints & Permissions  View PDF  View EPUB

ABSTRACT

Van Atteveldt et al. (2021)

Dictionary vs. supervised machine learning

- Supervised machine learning requires (potentially larger amounts) labeled data
- If the training sample is large enough supervised learning will outperform dictionaries

Additional considerations

- Hyperparameter selection
 - Via systematic comparison of different hyperparameters per algorithm
- Random undersampling (Galar et al., 2011)
 - Method to deal with unbalanced classes: use the max. number of positive instances per class and randomly sample the same number of instances of the negative class

Naïve Bayes Classifier

- Probabilistic classifier
 - Simple
 - Fast
 - Good Accuracy
-

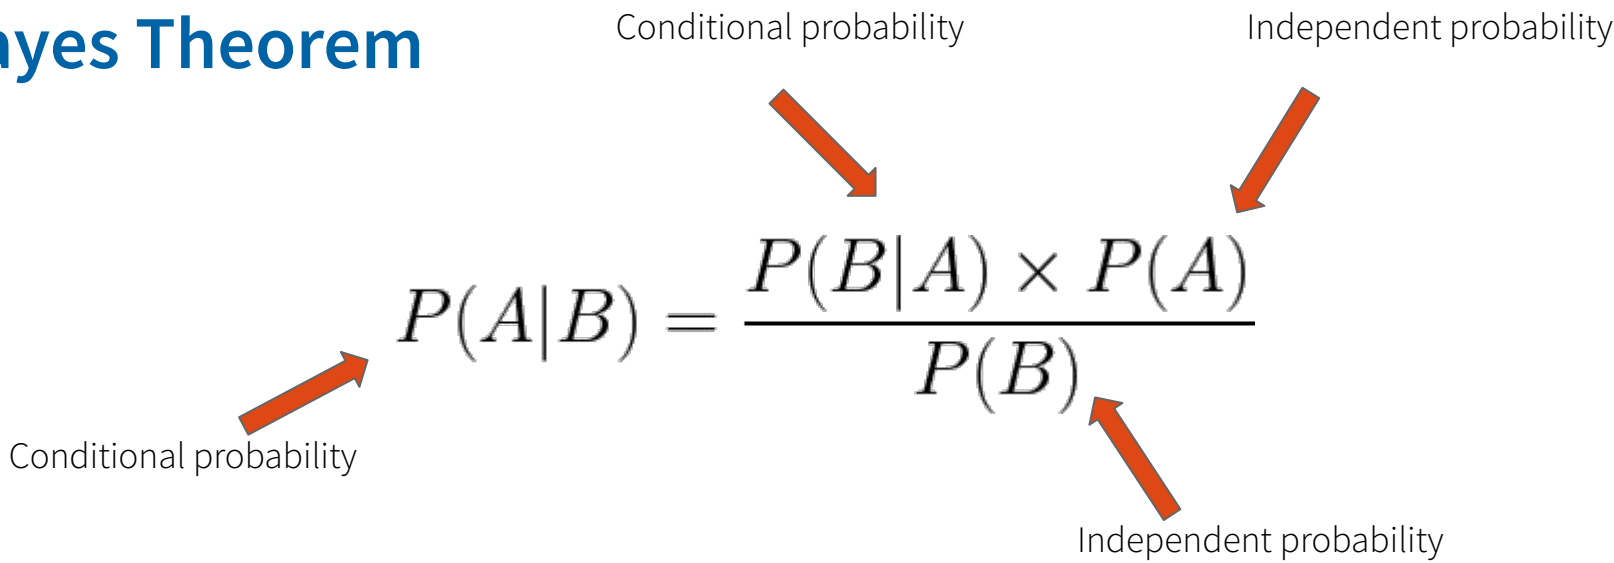
Bayes Theorem

$$P(A|B) = \frac{P(B|A) \times P(A)}{P(B)}$$

Bayes Theorem

Conditional probability

Independent probability


$$P(A|B) = \frac{P(B|A) \times P(A)}{P(B)}$$

Conditional probability

Independent probability

Bayes Theorem

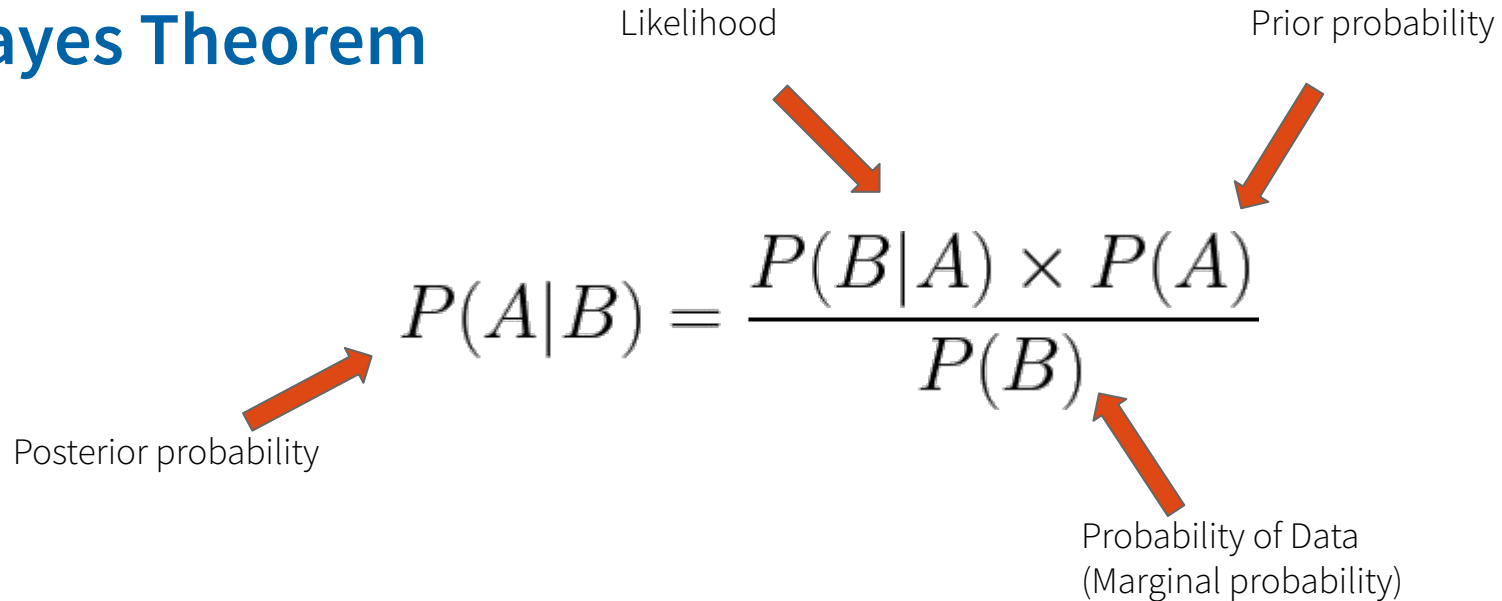


Diagram illustrating Bayes Theorem with labels and arrows:

- Likelihood** points to $P(B|A)$
- Prior probability** points to $P(A)$
- Posterior probability** points to $P(A|B)$
- Probability of Data (Marginal probability)** points to $P(B)$

$$P(A|B) = \frac{P(B|A) \times P(A)}{P(B)}$$

Bayes Theorem

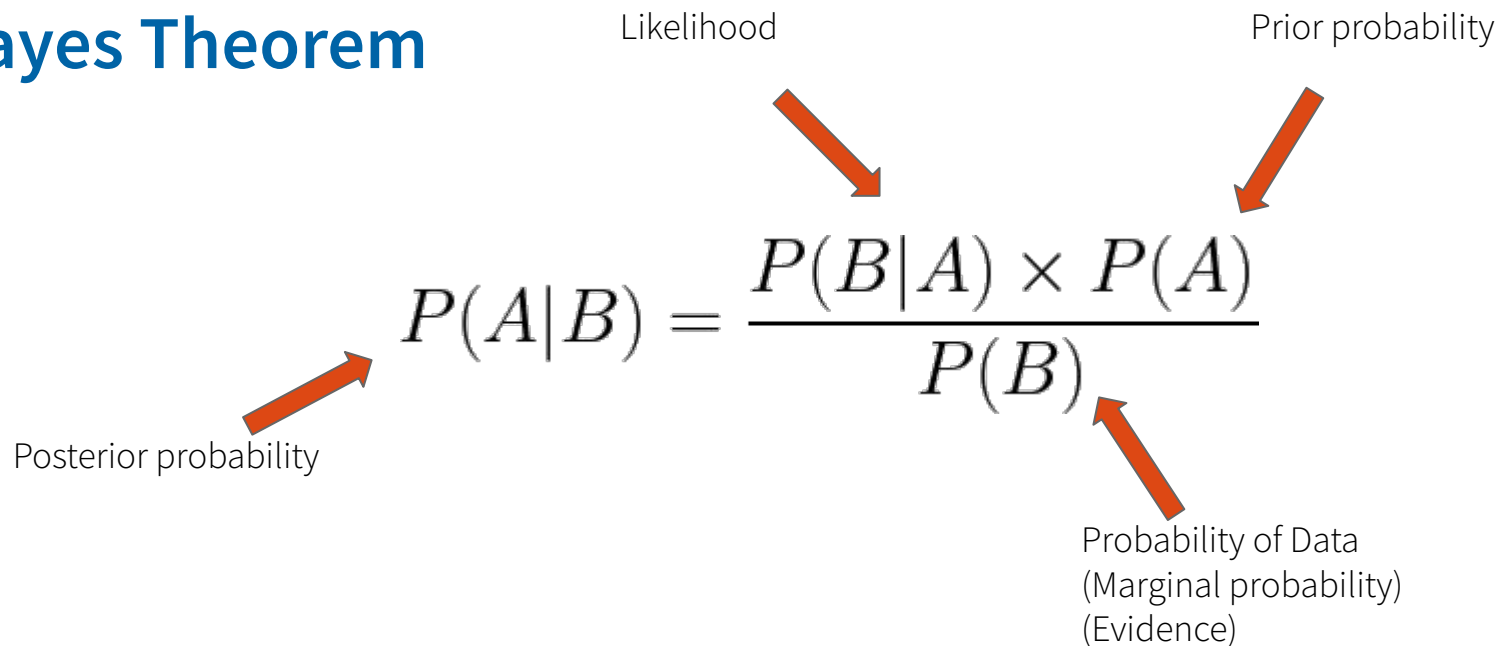
Likelihood

Prior probability

$$P(A|B) = \frac{P(B|A) \times P(A)}{P(B)}$$

Posterior probability

Probability of Data
(Marginal probability)
(Evidence)



The diagram illustrates the components of Bayes' Theorem. Four red arrows point from descriptive labels to specific parts of the equation $P(A|B) = \frac{P(B|A) \times P(A)}{P(B)}$. The label 'Likelihood' points to $P(B|A)$, 'Prior probability' points to $P(A)$, 'Posterior probability' points to $P(A|B)$, and 'Probability of Data (Marginal probability) (Evidence)' points to $P(B)$.

Bayes Theorem


$$P(A|B) \propto P(B|A) \times P(A)$$

Naïve Bayes Classifier

$$P(C_k | x_1, x_2, \dots, x_n)$$


Naïve Bayes Classifier

Class Features


$$P(C_k | x_1, x_2, \dots, x_n)$$

Naïve Bayes Classifier

Class Features


$$P(C_k | x_1, x_2, \dots, x_n)$$

Features are assumed to be independent. Hence, “**Naïve**”

Naïve Bayes Classifier

$$P(C_k|\mathbf{x}) = \frac{P(C_k) \times P(\mathbf{x}|C_k)}{P(\mathbf{x})}$$

Naïve Bayes Classifier

$$P(C_k|\mathbf{x}) \propto P(C_k) \times P(\mathbf{x}|C_k)$$

Naïve Bayes Classifier

$$\begin{aligned} p(C_k \mid x_1, \dots, x_n) &\propto p(C_k, x_1, \dots, x_n) \\ &\propto p(C_k) p(x_1 \mid C_k) p(x_2 \mid C_k) p(x_3 \mid C_k) \cdots \\ &\propto p(C_k) \prod_{i=1}^n p(x_i \mid C_k), \end{aligned}$$

Decision Rule

$$\hat{y} = \operatorname{argmax} p(C_k) \prod_{i=1}^n p(x_i|C_k)$$



Implemented in many stats/ML packages



Support Vector Machine

- Comes from computer science
 - Very good
 - Rather difficult math
-
- Considered one of the best of-the-shelf classification algorithms
-

Hyperplane

- $n-1$ dimensional plane that separates the n -dimensional space
-

Hyperplane

- n-1 dimensional plane that separates the n-dimensional space
- 2-dimensional hyperplane:

$$\beta_0 + \beta_1 X_1 + \beta_2 X_2 = 0$$

Hyperplane

- n-1 dimensional plane that separates the n-dimensional space
- 2-dimensional hyperplane:
- Line equation

$$\beta_0 + \beta_1 X_1 + \beta_2 X_2 = 0$$

Hyperplane

- n-1 dimensional plane that separates the n-dimensional space
- 2-dimensional hyperplane:
- Line equation

$$\beta_0 + \beta_1 X_1 + \beta_2 X_2 = 0$$



$$\beta_0 + \beta_1 X_1 + \beta_2 X_2 + \dots + \beta_p X_p = 0$$


Classification



$$\beta_0 + \beta_1 X_1 + \beta_2 X_2 + \dots + \beta_p X_p > 0$$

$$\beta_0 + \beta_1 X_1 + \beta_2 X_2 + \dots + \beta_p X_p < 0$$


Following images from:







SV Classifier



Support Vector Machine


- Non-linear version of the Support Vector Classifier
 - Extension using **Kernels**
-

$$f(x) = \beta_0 + \sum_{i \in \mathcal{S}} \alpha_i \langle x, x_i \rangle$$

$$f(x) = \beta_0 + \sum_{i \in \mathcal{S}} \alpha_i K(x, x_i)$$

$$f(x) = \beta_0 + \sum_{i \in \mathcal{S}} \alpha_i \langle x, x_i \rangle$$

$$f(x) = \beta_0 + \sum_{i \in \mathcal{S}} \alpha_i K(x, x_i)$$

 Kernel function



$$f(x) = \beta_0 + \sum_{i \in \mathcal{S}} \alpha_i \langle x, x_i \rangle$$

Kernel function

$$f(x) = \beta_0 + \sum_{i \in \mathcal{S}} \alpha_i K(x, x_i)$$

Polynomial Kernel
Non-linear

$$K(x_i, x_{i'}) = \left(1 + \sum_{j=1}^p x_{ij} x_{i'j}\right)^d$$




Kernel Trick

- ***Actual name***

Kernel Trick

- ***Actual name***
 - Attempt to place n -dimensional data into $n+1$ dimensional space
-






-5.0



-2.5

0.0
x1

2.5

5.0





Performing supervised machine learning in R and Python

R: quanteda, caret, e1071, klaR, C5.0, OneR

Python: scikit-learn

Coding session

- Sentiment
in movie
reviews



I liked it, but thought the third act nearly cratered the whole thing.

January 3, 2024 | [Full Review...](#)

Coding session

- Sentiment
in movie
reviews



I liked it, but thought the third act nearly cratered the whole thing.

January 3, 2024 | [Full Review...](#)



An intense, inventively filmed, and well-acted biopicture about the kinds of events that are hard on the heart.

[Full Review](#) | Original Score: 5/5 | Nov 20, 2023



In a movie about impending global catastrophe, he gives a close-up of a face, and a twitch of a lip the power of an atom bomb.

[Full Review](#) | Original Score: A | Nov 17, 2023



I liked it, but thought the third act nearly cratered the whole thing.

January 3, 2024 | [Full Review...](#)



An intense, inventively filmed, and well-acted biopicture about the kinds of events that are hard on the heart.

[Full Review](#) | Original Score: 5/5 | Nov 20, 2023

In a movie about impending global catastrophe, he gives a close-up of a face, and a twitch of a lip the power of an atom bomb.

[Full Review](#) | Original Score: A | Nov 17, 2023

Unsupervised classification

Day 3 Session 2

What do you see here?

- coffee
- cafe
- espresso
- latte
- barista
- beans
- brew
- cappuccino
- aroma
- roast

- programming
- code
- software
- development
- algorithm
- python
- function
- variable
- debugging
- Java

- fitness
 - exercise
 - health
 - workout
 - gym
 - nutrition
 - weight
 - strength
 - cardio
 - yoga
-

Topic 1: Label?

- coffee
- cafe
- espresso
- latte
- barista
- beans
- brew
- cappuccino
- aroma
- roast

Topic 2: Label?

- programming
- code
- software
- development
- algorithm
- python
- function
- variable
- debugging
- Java

Topic 3: Label?

- fitness
 - exercise
 - health
 - workout
 - gym
 - nutrition
 - weight
 - strength
 - cardio
 - yoga
-

Types of machine learning

1. Supervised
 - An outcome variable is defined
 - Focus is on prediction
 2. Unsupervised
 - No outcome variable has been defined
 - Focus is on patterns
-

How to use the *supervised* methods?

- Easy
 - At least *conceptually*
 - Clear **objective function**
-

How to use the *supervised* methods?

$$Y = (y_1, y_2, \dots, y_n)$$

$$X = (x_1, x_2, \dots, x_n)$$

$$(y_1, x_1), (y_2, x_2), \dots, (y_n, x_n)$$

Task to predict \hat{y} as close to y

How to use the *supervised* methods?

$$L(y, \hat{y}) = (y - \hat{y})^2$$

$$\hat{y} = \operatorname{argmin}_{\theta} E [L(model(\mathbf{x}, \theta), y)]$$

How to use unsupervised learning

- Objective function?

How to use unsupervised learning

- Objective function?
 - Quantity of interest?
-

How to use unsupervised learning

- Objective function = ***your*** quantity of interest



How to use unsupervised learning

- Objective function = ***your*** quantity of interest
 - This is difficult
-

Measurement

- A collection of quantitative or numerical data that describes a property of an object or event

Measurement

- A collection of quantitative or numerical data that describes a property of an object or event
 - What is the ***object***?
-

Measurement (in the social sciences)

- Operationalisation \Rightarrow Data Collection \Rightarrow Analyses \Rightarrow Measurement
 - The quantity of interest is **extrinsic** to the model
-

Measurement (in computer science)

- Model Building \Rightarrow Measurement of Performance
 - The quantity of interest is ***intrinsic*** to the model
-

$$L(\theta, \hat{\theta}) = (\hat{\theta} - \theta)^2$$

$$L(\theta, \hat{\theta}) = (\hat{\theta} - \theta)^2$$

Main quantity of Interest for
computer science

$$L(\theta, \hat{\theta}) = (\hat{\theta} - \theta)^2$$

Means to an end
for social science

Main quantity of Interest for
computer science

$$L(\theta, \hat{\theta}) = (\hat{\theta} - \theta)^2$$

Means to an end
for social science

Main quantity of Interest for
computer science

$$\hat{\theta} \approx \theta$$

What social science wants

Prediction vs. Inference

- Computer scientists often emphasise ***prediction***
 - Social scientists are often more interested in ***inference***
 - Vast, multidimensional parameter space = not suitable for inference
 - Good for prediction
 - E.g., Turing test
 - Machine passes
 - Why does it pass/not pass
-

Problems

- Translation of Social Science concepts
 - Connecting Methods to Theory
 - Difficult to understand what is being ***measured***
-

Unsupervised Learning Example

- Clustering algorithms are great tools
 - Not well suitable for the “standard” social science paradigm
 - Needs external validation, but there is no “best” method
 - “Validation” based on “theory” or “expectation” leads to biases
-

The paradigm

- Approximating a data-generating process

The paradigm

- Approximating a data-generating process
 - Assumption: there is one (and only one) “true” data-generating process
 - It **is** the reality
-

Sticking to the paradigm

- The “normal” paradigm works only if we assume that there is one “correct” classification
 - Need to adapt to different methods
-

Sticking to the paradigm

- The “normal” paradigm works only if we assume that there is one “correct” classification
 - Need to adapt to different methods
 - Unsupervised methods are **meaningless** in conjunction with the “true” data-generating process assumption
-

Focus is on Discovery

Objectives

Descriptive analysis/Discriminating words:

- What are the characteristics of a corpus? How do some documents compare to each other
- Keyness, collocation analysis, readability scores, Cosine/Jaccard similarity

Clustering and scaling:

- What groups of documents are in the corpus? Can the documents be placed on a dimension?
- Cluster analysis, principal component analysis, wordfish..

Topic modeling:

- What are the main themes in a corpus?
- LDA, STM, BERTopic

K-Means Clustering

- Simple(ish) algorithmic method
 - Partitions the data into K non-overlapping clusters
-

Setup

$$C_1, C_2, \dots, C_K$$

$$C_1 \cup C_2 \cup \dots \cup C_K = \{1, \dots, n\}$$

$$C_k \cap C_{k'} = \emptyset \text{ for all } k \neq k'$$

Assumption and task

- Optimal clustering solution is the one where ***within-cluster variation*** is ***as small as possible***

$$W(C_k)$$

$$\text{minimize}_{C_1, \dots, C_K} \left\{ \sum_{k=1}^K W(C_k) \right\}$$

Within cluster variation

$$W(C_k) = \frac{1}{|C_k|} \sum_{i, i' \in C_k} \sum_{j=1}^p (x_{ij} - x_{i'j})^2$$

$$\text{minimize}_{C_1, \dots, C_K} \left\{ \sum_{k=1}^K \frac{1}{|C_k|} \sum_{i, i' \in C_k} \sum_{j=1}^p (x_{ij} - x_{i'j})^2 \right\}$$

Algorithm

1. Randomly assign cluster numbers (1 through K) for each observation
 2. Iterate until no further changes to the cluster assignment:
 - a. For each cluster determine the **centroid** (average of all observations in the cluster)
 - b. Re-assign observations to a cluster with the closest centroid (calculated with a distance metric).
-

Guarantees convergence at a *local optimum*

- Cannot guarantee the best solution
 - But are rather good one
 - Sensitive to random assignment at the start
-

Cluster Algorithms Validation

- Data assumptions (think data generation)
 - Internal validity (best results for the data)
 - External validity (matches with pre-existing understanding of data)
 - Cross-validity (similar results across similar datasets)
 - ***You are the validation method***
-

Topic Modelling

- A model to discover latent topics
 - **Not** synonymous with LDA
 - LDA is one of topic models
-
- Latent semantic analysis
 - Singular value decomposition
 - Even clustering methods (like the one we just discussed)
-

Latent Dirichlet Allocation

- Bayesian generative hierarchical model
 - First introduced as a way to simultaneously model traits and genes (Pritchard, 2000)
 - Adjusted for text analysis ML applications (Blei et al., 2003)
-

Latent Dirichlet Allocation

- Estimates a distribution of **words** across **documents** across **latent topics**
-

Latent Dirichlet Allocation

- By modeling distributions of topics over words and words over documents, topic models identify the most discriminatory groups of documents automatically.

Assumption: if a document is about a certain topic, one would expect words, that are related to that topic, to appear in the document more often than in documents that deal with other topics.

Mixture Models

$$P(x) = \sum_{k=1}^K \pi_k \times P(x | \theta_k)$$

- $P(x)$ represents the probability density function (PDF) of the observed data point x
 - K is the number of components in the mixture model
 - π_k is the mixing proportion or weight assigned to the k -th component, representing the probability of a data point belonging to that component. These weights satisfy the condition $\sum(\pi_k) = 1$ and $0 \leq \pi_k \leq 1$
 - $P(x|\theta_k)$ represents the conditional probability of observing data point x given the parameters θ_k of the k -th component distribution
-

Hierarchical Models

Hierarchical Models

$$y = \beta_0 + \beta_1 X_1 + \dots + \beta_k X_k + \epsilon$$

Hierarchical Models

$$y \sim \text{Normal}(\mu, \sigma)$$

$$\mu = \beta_0 + \beta_1 X_1 + \dots + \beta_k X_k$$

$$\beta \sim \text{Normal}(0, 5)$$

$$\sigma \sim \text{Exponential}(1)$$

Hierarchical Models

Likelihood Function

$$y \sim \text{Normal}(\mu, \sigma)$$

$$\mu = \beta_0 + \beta_1 X_1 + \dots + \beta_k X_k$$

$$\beta \sim \text{Normal}(0, 5)$$

$$\sigma \sim \text{Exponential}(1)$$

Prior distribution of the parameters

Hierarchical Models

$$y \sim \text{Normal}(\mu, \sigma)$$

$$\mu = \beta_0 + \beta_1 X_1 + \dots + \beta_k X_k$$

$$\beta \sim \text{Normal}(0, 5)$$

$$\sigma \sim \text{Exponential}(1)$$

Hyperparameters

LDA

$$\theta_k \sim \textit{Dirichlet}(\alpha)$$

for each topic k

$$\eta_d \sim \textit{Dirichlet}(\beta)$$

for each document d

$$z_{d,w} \sim \textit{Multinomial}(\eta_d)$$

for each word w in document d

$$x_{d,w} \sim \textit{Multinomial}(\theta_{z_{d,w}})$$

for each word w in document d

LDA





- θ_k is the topic-word distribution for topic k , representing the probability of each word given the topic.
 - η_d is the document-topic distribution for document d , representing the probability of each topic in the document.
 - $z_{d,w}$ is the topic assignment for word w in document d , indicating which topic generated the word.
 - $x_{d,w}$ is the observed word in document d .
-

LDA

- For each topic $k \in \{1, \dots, K\}$:
 - Draw a distribution over words $\theta_k \sim \text{Dirichlet}(\alpha)$, where α is a hyperparameter representing the topic-word prior.
 - For each document $d \in \{1, \dots, D\}$:
 - Draw a distribution over topics $\eta_d \sim \text{Dirichlet}(\beta)$, where β is a hyperparameter representing the document-topic prior.
 - For each word w in document d :
 - Draw a topic assignment $z_{d,w} \sim \text{Multinomial}(\eta_d)$, indicating which topic generated the word.
 - Draw a word $x_{d,w} \sim \text{Multinomial}(\theta_{z_{d,w}})$, indicating the specific word generated by the chosen topic.
-

LDA

- The goal of LDA is to infer the posterior distributions of the latent variables θ and η given the observed documents.
 - Once the posterior distributions are estimated, LDA can be used to assign topics to new documents or extract the most probable words for each topic.
-



Tricky to estimate

$$\begin{aligned}
& \int_{\theta_j} P(\theta_j; \alpha) \prod_{t=1}^N P(Z_{j,t} \mid \theta_j) d\theta_j = \int_{\theta_j} \frac{\Gamma\left(\sum_{i=1}^K \alpha_i\right)}{\prod_{i=1}^K \Gamma(\alpha_i)} \prod_{i=1}^K \theta_{j,i}^{n_{j,(\cdot)}^i + \alpha_i - 1} d\theta_j \\
&= \frac{\Gamma\left(\sum_{i=1}^K \alpha_i\right)}{\prod_{i=1}^K \Gamma(\alpha_i)} \frac{\prod_{i=1}^K \Gamma(n_{j,(\cdot)}^i + \alpha_i)}{\Gamma\left(\sum_{i=1}^K n_{j,(\cdot)}^i + \alpha_i\right)} \int_{\theta_j} \frac{\Gamma\left(\sum_{i=1}^K n_{j,(\cdot)}^i + \alpha_i\right)}{\prod_{i=1}^K \Gamma(n_{j,(\cdot)}^i + \alpha_i)} \prod_{i=1}^K \theta_{j,i}^{n_{j,(\cdot)}^i + \alpha_i - 1} d\theta_j \\
&= \frac{\Gamma\left(\sum_{i=1}^K \alpha_i\right)}{\prod_{i=1}^K \Gamma(\alpha_i)} \frac{\prod_{i=1}^K \Gamma(n_{j,(\cdot)}^i + \alpha_i)}{\Gamma\left(\sum_{i=1}^K n_{j,(\cdot)}^i + \alpha_i\right)}.
\end{aligned}$$

$$\begin{aligned}
& \int_{\varphi} \prod_{i=1}^K P(\varphi_i; \beta) \prod_{j=1}^M \prod_{t=1}^N P(W_{j,t} \mid \varphi_{Z_{j,t}}) d\varphi \\
&= \prod_{i=1}^K \int_{\varphi_i} P(\varphi_i; \beta) \prod_{j=1}^M \prod_{t=1}^N P(W_{j,t} \mid \varphi_{Z_{j,t}}) d\varphi_i \\
&= \prod_{i=1}^K \int_{\varphi_i} \frac{\Gamma\left(\sum_{r=1}^V \beta_r\right)}{\prod_{r=1}^V \Gamma(\beta_r)} \prod_{r=1}^V \varphi_{i,r}^{\beta_r-1} \prod_{r=1}^V \varphi_{i,r}^{n_{(\cdot),r}^i} d\varphi_i \\
&= \prod_{i=1}^K \int_{\varphi_i} \frac{\Gamma\left(\sum_{r=1}^V \beta_r\right)}{\prod_{r=1}^V \Gamma(\beta_r)} \prod_{r=1}^V \varphi_{i,r}^{n_{(\cdot),r}^i + \beta_r - 1} d\varphi_i \\
&= \prod_{i=1}^K \frac{\Gamma\left(\sum_{r=1}^V \beta_r\right)}{\prod_{r=1}^V \Gamma(\beta_r)} \frac{\prod_{r=1}^V \Gamma(n_{(\cdot),r}^i + \beta_r)}{\Gamma\left(\sum_{r=1}^V n_{(\cdot),r}^i + \beta_r\right)}.
\end{aligned}$$
